

DICT Acceptable Use Policy for AI Productivity Tools  
Document ID: DICT-AUP-2026-01  
Effective Date: 2026-04-01  
Review Cycle: Annual  
Owner: DICT Information Security Office

## 1. Purpose

This policy defines the acceptable use of AI productivity tools, including cloud-hosted assistants and locally deployed open-weight models, by DICT personnel and accredited partners. It applies to every device managed by DICT and to every account that accesses DICT data.

## 2. Scope

2.1 In scope: any system or service that processes, summarises, drafts, or generates text, code, images, or audio using a machine learning model.

2.2 Out of scope: deterministic tooling such as spreadsheets, version control, and ticketing systems unless they embed an AI feature.

2.3 Out of scope: offline spell checkers and grammar checkers that do not send data to a remote service.

## 3. Permitted uses

3.1 Drafting internal memos, briefings, and meeting summaries from source material the user is authorised to read.

3.2 Generating, refactoring, and reviewing source code for DICT projects, provided the source code is not classified as Restricted or higher.

3.3 Translating documents between English and Filipino for official communications.

3.4 Producing first-pass design mockups from a written requirements document.

## 4. Prohibited uses

4.1 Submitting data classified as Restricted or higher to any cloud-hosted AI service. Restricted data includes citizen personal data, internal audit reports, and infrastructure credentials.

4.2 Generating, distributing, or storing content that is unlawful, discriminatory, harassing, or fraudulent.

4.3 Using AI output as the sole basis for a decision that affects a citizen's rights, benefits, or obligations without human review.

4.4 Bypassing the human-in-the-loop controls specified in Section 6 for any autonomous workflow.

## 5. Model selection

5.1 Cloud models (Gemini, OpenAI, Anthropic) may be used for tasks that do not involve Restricted data.

5.2 Locally deployed open-weight models (Gemma 4, Llama 3) are preferred for tasks that involve internal or sensitive-but-unclassified data and for work performed in air-gapped environments.

5.3 The Information Security Office maintains a list of approved models and their permitted data classes.

## 6. Human oversight

6.1 Every AI-generated artefact that informs an external communication must be reviewed by a human before it is sent.

6.2 Autonomous workflows require a named human approver, an audit log entry, and a tested rollback path before they are activated.

6.3 The audit log must record the prompt, the model, the output, the decision, and the action with timestamps.

## 7. Incident reporting

7.1 Suspected policy violations must be reported to the Information Security Office within 24 hours.

7.2 Confirmed violations may result in revocation of AI tool access, referral to the DICT disciplinary process, or both.

## 8. Review

This policy is reviewed annually by the Information Security Office in consultation with the DICT Legal Division. Material changes are approved by the DICT Secretary before they take effect.